

Professional Summary

Performance-oriented Data Engineer with 4+ years of experience in system performance optimization, distributed data processing, and cloud-native development. Skilled in Spark tuning, profiling, and large-scale ETL optimization across Azure, AWS, and GCP. Experienced with Kafka-based event streaming, Kubernetes deployments, and cloud monitoring tools for performance analysis and reliability.

Professional Experience

Honeywell – Data Engineer | Feb 2021 – Jan 2024

- Designed and optimized PySpark ETL pipelines on Azure Databricks, migrating 5+ TB of data with up to 40% faster performance.
 - Performed Spark job profiling and cluster tuning to reduce runtime and resource costs by 30%.
 - Implemented monitoring and performance tracking using Databricks dashboards and Azure Monitor.
 - Built and automated ADF pipelines integrating MySQL, Hive, and ADLS with dependency-driven scheduling.
 - Integrated MongoDB with Azure Databricks for data enrichment of semi-structured datasets, enhancing query performance and analytics.
 - Applied RBAC and Key Vault security controls, ensuring compliant access management and data protection.
 - Collaborated with QA and DB teams to identify and resolve system-level performance bottlenecks.
-

Cognizant Technology Solutions – Software Engineer | Sep 2019 – Dec 2020

- Built and optimized PySpark ETL jobs on AWS EMR, processing over 1 TB/week of enterprise data.
 - Developed NiFi workflows for automated ingestion from MySQL to Hive/S3, enabling data flow.
 - Integrated Kafka streams with Spark for real-time data ingestion and event-driven architecture.
 - Tuned Spark transformations using CloudWatch metrics, improving performance by 25%.
 - Developed SQL-based validation scripts to ensure 99.5% data accuracy in production pipelines.
-

Education

University of Massachusetts – Dartmouth, MA

M.S. in Data Science (GPA: 4/4)

Technical Skills

- **Performance & Monitoring:** Spark tuning, job profiling, Databricks performance dashboards, Azure Monitor, AWS CloudWatch
- **Cloud Platforms:** Azure (Databricks, ADF), AWS (EMR, Glue, S3), GCP (Dataproc, BigQuery)
- **Streaming & Distributed Systems:** Apache Kafka, NiFi, Flink
- **Cloud-Native & DevOps:** Docker, Kubernetes, Helm (basic), Jenkins, GitHub Actions
- **Programming & Databases:** Python, Java, SQL, PL/SQL, Oracle, MongoDB
- **Security & Compliance:** IAM, RBAC, Azure Key Vault, encryption, awareness of HIPAA/FHIR standards