Praveen Nemalikanti

Azure Data Engineer

Phone: 6092480042
Email: praveennemalikanti91@gmail.com
LinkedIn: https://www.linkedin.com/in/praveen-nemalikanti-1788513aa/

PROFESSIONAL SUMMARY

- With over 10 + years of comprehensive experience in Data Engineering, I have excelled in architecting and implementing scalable **Azure, Snowflake, Big Data, and ETL** solutions. My expertise optimizes performance and drives insights across diverse domains including Banking, Finance, Healthcare, Insurance, and Education.
- Proficient in **Azure services** like **Azure Data Factory, Azure Databricks, Azure Data Lake Storage Gen2, Azure Logic Apps, Azure Functions, Azure Event Hubs, Azure Synapse Analytics, Azure SQL Database, Azure Cosmos DB, Azure Stream Analytics, Azure Blob Storage, Azure Purview Azure Analysis Services, Azure Machine Learning, Azure Monitor, Azure Kubernetes Service (AKS), Azure DevOps, Azure Cognitive Services.**
- Designed and implemented roles and access privileges in **Azure Synapse Analytics, Azure SQL Database, and Azure Data Lake Storage (ADLS) Gen2**, ensuring robust data security protocols and reducing unauthorized access incidents.
- Integrated **Azure Data Lake Storage (ADLS)** Gen2 with **Azure Synapse Analytics** for distributed data processing, utilizing **Informatica PowerCenter** and **SQL Server Integration Services (SSIS)** for Actions, Transformations, and ETL tasks.
- Designed scalable Azure data solutions, optimizing models and integrating Synapse Analytics with **Unity Catalog, Data Catalog, Databricks,** and **Event Hubs** for governance, metadata, processing, and streaming, ensuring impactful outcomes
- Orchestrated automated data pipelines on **Azure Databricks** and **Azure Data Factory (ADF)**, using **Azure Synapse Analytics** for cloud-based data warehousing. Executed seamless integration pipelines between Azure Synapse, **Azure Purview,** and other Azure services, leveraging Purview for comprehensive data governance and lineage insights.
- Implemented **Medallion** architecture in Azure, utilizing Azure Functions for event-driven processing, **Azure Data Lake Storage (ADLS) Gen2** for scalable data storage, and **Azure Synapse Analytics** for real-time analytics. Ensured efficient, scalable data processing workflows, leveraging **Azure Purview** for comprehensive data governance.
- Hands-on experience handling various file formats (**CSV, JSON, Parquet, Avro, etc**.) to accommodate diverse data sources and optimize data storage efficiency. Proficient in Informatica, **SSIS (SQL Server Integration Services)**, and **SSRS (SQL Server Reporting Services)** for comprehensive data integration and reporting solutions.
- Hands-on experience in legacy data migration from **Teradata** to **Azure Synapse Analytics** and on-premises to Azure Cloud, leveraging **Azure Functions** for serverless applications. Implemented **Azure Data Factory** to orchestrate complex data workflows seamlessly between on-premises and Azure environments.
- Orchestrated data pipelines across **Azure Databricks** and Glue, utilizing Azure Synapse Analytics, Redshift, **Azure Logic Apps**, and **Airflow (Directed Acyclic Graphs) DAGs** for seamless cloud-based data warehousing.
- Implemented robust ETL workflows, ensuring governance and data quality, utilizing Azure databases such as **Azure SQL Database**, **Azure Synapse Analytics, Cosmos DB, and Azure Cache**. Integrated **Azure Fabric** for orchestrating data workflows, enhancing scalability and efficiency across Azure services.
- Successfully migrated legacy data from on-premises and multi-cloud environments to **Azure Synapse Analytics and Redshift**, employing serverless architectures with **Azure Functions and Lambda**.
- Designed and enforced role-based access control (**RBAC**) policies across Azure Synapse Analytics, integrated **Azure Key Vault** ensuring comprehensive data security and encryption key management in multi-cloud environments.
- Utilized Azure Synapse Analytics for optimized querying with clustering, partitioning, and materialized views. Integrated **Kafka, Delta Tables, Hive, Databricks, PySpark, DataFrames, RDDs**, and JSON to enhance real-time streaming and data processing, leveraging **DAG**s and **API**s for seamless interaction across analytics workflows.

- Implemented sophisticated partitioning strategies within **Snowflake** to significantly enhance query performance for critical analytics tasks, incorporating **SCD** types for effective management of historical data changes
- Integrating data storage solutions within **Spark**, focusing on seamless integration **with Azure Data Lake Storage and Blob Snowflake Storage**, leveraging **Azure Fabric** for enhanced scalability and reliability.
- Collaborated with stakeholders in the Financial and healthcare domains to architect and implement data solutions on cloud platforms like **Snowflake and Azure Databricks**, designing scalable and robust data architectures to support real-time analytics, reporting, and decision-making processes within the financial, and educational industries.
- Utilized **Snowpipe** for real-time data ingestion in Snowflake, automating data loading processes and minimizing manual intervention. Implemented **SCD (Slowly Changing Dimension)** types 1 and 2 for effective management of historical data.
- Conducted performance tuning of **Big Data** clusters, designed end-to-end data pipelines, and developed intricate database solutions using **Teradata, Oracle, SQL Server, PLSQL, and PostgreSQL.**
- Designed and maintained ETL processes, integrating diverse data sources into databases and warehouses with **Unity Catalog** for governance, **Azure Data Catalog** for metadata, and **Medallion** architecture for organization.
- Developed and maintained scalable data processing pipelines using **Python** libraries like **Pandas, NumPy, scikit-learn**, **Apache Spark, and Java**. Ensured efficient, reliable, and accurate data transformations for large-scale datasets.
- Utilized Azure Data Factory for orchestration, **Azure Event Grid** for real-time event processing, and **Azure DevOps** for **CI/CD** in data pipeline deployment. Integrated Azure Functions for efficient serverless computing in event-driven cases.
- Extensive proficiency in **Big Data technologies, Hadoop Ecosystem, Spark**, and various **NoSQL** databases, actively enhancing query performance and facilitating smooth data transmission.
- Leveraged **Apache Spark** for complex data transformations, implemented partitioning and optimization strategies in **Hive** and configured **Zookeeper** for cluster coordination. Utilized **RDD (Resilient Distributed Dataset)** for efficient distributed data processing and fault tolerance within Spark jobs.
- Partnered with DevOps teams, leveraging **Azure DevOps** for **CI/CD** in data pipeline deployment, and collaborated with **Jenkins** for comprehensive end-to-end automation, including database implementation.
- Collaborated with cross-functional teams using **Agile** and **Scrum** methodologies to deliver organized data solutions. Ensured data quality, governance, and optimized infrastructure to meet stakeholders' needs, empowering data analysts.

TECHNICAL SKILLS

Cloud & Data Platforms	Azure Data Factory, Azure Databricks, Azure Synapse Analytics, Azure Data Lake Storage Gen2, Azure Fabric, Azure Purview, Snowflake, Azure SQL Database, Azure Blob Storage
Big Data & Streaming	Apache Spark, PySpark, Spark SQL, Delta Lake, Kafka, Event Hubs, Hive
Hadoop Distribution	Cloudera, Horton Works
Languages	Python, Java, SQL, PL/SQL, HiveQL, Scala, T-SQL, PySpark
Operating Systems	Windows (XP/7/8/10), UNIX, LINUX, UBUNTU, CENTOS
Databases	Teradata, Oracle DB, PostgreSQL, SQL server, MSSQL, MS Access
Scheduling	IBM Tivoli, Control-M, Oozie, Airflow
Version Control	GIT, GitHub, Bitbucket, VSS
Methodology	Agile, Scrum
IDE &Build Tools, Design	Eclipse, PyCharm IDE, Visual Studio
Analysis and Visualization	Power BI, Tableau, QuickSight

EDUCATION:

- **Master's in Computer Science & Info Tech at Sacred Heart University, USA – December, 2014**
- **Bachelor's in Electronics and Communication Engineering, Vignan University, India – May, 2013**

CERTIFICATIONS:

- **Microsoft Certified: Azure Data Engineer Associate**
- **SnowPro Core Certification**
- **Databricks Certified Data Engineer Professional**
- **AWS Certified: Cloud Practitioner**
- **AWS Certified: Solutions Architect Associate**

<u>WORK EXPERIENCE</u>
Role: Senior Data Engineer | Aug 2023 – Till Now
Client: Penn Medicine, Philadelphia, PA
Responsibilities:

- Developed impactful ETL workflows using Azure Data Factory to integrate diverse, streaming financial data into Azure Synapse Analytics for fraud detection. Leveraged Kafka, Azure Event Hubs, Delta Tables, Hive, Databricks, PySpark, Data Frames, Spark SQL, RDDs, DAGs, JSON, SSIS, Informatica, and IICS to ensure high-quality, reliable analytics.
- Utilized JSON, Avro, Parquet, and ORC data formats in Azure Data Lake Storage (ADLS) Gen2 for 25% improved storage efficiency. Integrated Azure Data Factory (ADF) for structured data modeling, enhancing data consistency across formats. Leveraged Azure Databricks and PySpark for efficient data processing, optimizing storage and query performance.
- Implemented Azure RBAC (Role-Based Access Control) and ACL (Access Control List) for Data Lake Storage Gen2 to enforce row-level and column-level security, ensuring precise control over data access based on roles and permissions.
- Initiated the setup of essential Azure Cloud services for scalability, flexibility, and cost optimization. Established a robust Azure infrastructure to support effective data engineering operations. Integrated Azure Databricks and PySpark, enhancing scalability, flexibility, and cost efficiency in data processing and analytics tasks within the Azure ecosystem.
- Improved Azure Synapse Analytics performance by implementing keys and compression strategies. Utilized Azure Databricks and Kafka for data modeling and streaming. Integrated Azure Fabric to orchestrate data workflows and enhance scalability across Azure services, incorporating APIs for streamlined integration management of data workflows.
- Managed the implementation of Azure Purview for comprehensive data governance, overseeing discovery, cataloging, and compliance across diverse environments. Maintained Purview's Data Catalog to annotate data assets with business glossaries and classifications, ensuring transparency, trust, and clear data lineage management.
- Using Terraform, integrated with various Azure services and resources commonly used in data engineering, such as Azure Blob Storage for object storage, Azure Synapse Analytics for data warehousing, Azure SQL Database for relational databases, Azure Data Factory (ADF) for ETL jobs, and Azure Functions for serverless computing.
- Integrated machine learning models into Azure Machine Learning, utilizing Azure Data Factory, Databricks, and Kafka for real-time data processing, which contributed to a $100,000 revenue increase through improved data insights.
- Designed and implemented robust data ingestion and storage strategies, leveraging Azure Blob Storage, Azure Synapse Analytics, and Azure Data Factory (ADF) to ensure efficient and scalable solutions. Integrated DBT for managing data transformations and modeling within the strategies, enhancing data consistency and usability.
- Integrated DBT into Terraform configurations for managing data transformations and modeling within the infrastructure as a code paradigm, leveraging Azure Databricks for scalable data processing. Utilized Apache Airflow to orchestrate workflows with DAGs (Directed Acyclic Graphs), ensuring efficient task scheduling and monitoring in the data pipeline.
- Demonstrated proficiency in SQL databases (MSSQL Server, MySQL, Oracle DB, PostgreSQL, MongoDB) and expertly formulated cloud migration strategies using Azure services, resulting in successful migration from On-Premises to Cloud.
- Orchestrated the seamless migration of databases to the cloud, leveraging Informatica, SSIS, and Azure services, resulting in a remarkable 30% reduction in on-premises infrastructure costs while maintaining data consistency and accessibility.
- Implemented data management and ingestion pipelines across bronze, silver, and gold layers: linked services for seamless data source connectivity in bronze, real-time streaming with self-hosted integration runtime in silver, and advanced analytics for sophisticated insights in gold. Developed APIs to facilitate integration and interaction in these data layers.
- Optimized Azure infrastructure with automation via Ansible playbooks, reducing manual intervention by 30%. Integrated DBT, Kafka, Unity Catalog, and Snowflake for efficient data and infrastructure management.
- Designed and implemented data models optimized for Slowly Changing Dimensions (SCD), ensuring historical accuracy and integrity through methodologies in Informatica, IICS, SSIS, Azure Data Factory, and Airflow.

- Designed and implemented end-to-end CI/CD pipelines using Jenkins for continuous integration and deployment. Achieved a 25% reduction in time-to-deployment, enhancing the speed and reliability of software releases.
- Implemented Schema RDDs with Spark SQL, reducing processing complexities by 20% in Azure Databricks. Utilized DBT and Kafka for data pipeline management, incorporating SCDs (Slowly Changing Dimensions) for efficient historical data management within the pipelines. Integrated Snowflake for enhanced data warehousing and scalable analytics.
- Took charge of setting up and building Azure infrastructure (VPC, Azure Blob Storage, Azure Cosmos DB, Azure IAM, Azure Disk Storage, Azure DNS, Azure Service Bus, Azure Monitor, Azure Security Center, Azure Virtual Machine Scale Sets, and Azure SQL Database) using Azure Resource Manager templates, ensuring a robust and scalable environment.
- Implemented best practices for data engineering on Azure Databricks, ensuring reliable, scalable, and high-performance data processing. Proficient in job workflow scheduling and locking tools like Azure Data Factory, Oozie, Zookeeper, Airflow, and Apache Nifi. Also developed APIs for seamless integration with data processing workflows.
- Orchestrated the design and implementation of Delta Lake alongside traditional Data Lake solutions for scalable data storage, optimizing processing pipelines for efficiency. Integrated SQL for enhanced querying and data manipulation capabilities within the Delta Lake architecture.
- Utilized PySpark for data manipulation and transformation tasks, scripting Azure interactions with Azure SDK for Python to seamlessly integrate with Delta Lake and Data Lake architectures. These efforts were instrumental in achieving performance, cost-effectiveness, and reliability goals in data workflows.
- Crafted compelling Power BI and Tableau dashboards for clients, leveraging insights from both Delta Lake and Data Lake repositories stored in Azure Data Lake Storage (ADLS) Gen2. Integrated SQL for advanced querying capabilities, resulting in a 20% enhancement in data accessibility and visualization effectiveness.
- Deployed Azure Synapse Analytics alongside Azure SQL Database for seamless ad-hoc data analysis and querying augmenting flexibility, and speed in extracting actionable insights from diverse data sources, including those stored in Azure Data Lake Storage (ADLS) Gen2 and Azure SQL Server databases.
- Integrated Azure Event Hubs and Azure Service Bus with Delta Lake, Data Lake, and Azure SQL Server for real-time processing and messaging within an Agile Scrum framework and onshore-offshore model, enhancing communication, responsiveness, and operational efficiency.
- Established robust configurations and optimizations for Azure Synapse Analytics, ensuring efficient and cost-effective ad-hoc querying capabilities on Azure Data Lake Storage (ADLS) Gen2, contributing to a 30% reduction in query response.
- Designed Terraform configurations to scale Azure infrastructure dynamically based on workload demands and performance requirements, optimizing resource utilization and costs. Additionally, integrated APIs for enhanced automation and orchestration capabilities.
- Implemented Kubernetes alongside Docker to facilitate auto-scaling and CI. Docker images were uploaded to the registry to enable deployment through Kubernetes. Utilized the Kubernetes dashboard for service monitoring and management.
- Defined data SLAs and implemented automated data quality checks across critical pipelines, reducing downstream data incidents and improving trust in analytics and reporting for business stakeholders.
- Implemented Azure Monitor for comprehensive monitoring and management of resources, setting up alerts, and collecting essential metrics, ensuring optimal system performance and reliability.
- Established scheduling and job automation using Azure Data Factory (ADF), Oozie, and Airflow, streamlining data processing workflows and achieving a 30% reduction in processing time. PySpark was utilized extensively within these workflows to handle large-scale data processing tasks efficiently.

Environment: Azure Data Factory (ADF), Azure Databricks, Azure Data Lake Storage (ADLS) Gen2, Azure Fabric, Azure Purview, Azure SQL Database, Azure Event Hubs, Azure Functions, Azure Synapse Analytics, Snowflake, Kafka, Talend, Snow SQL, Snowpipe, Change Data Capture (CDC), Snowflake Time Travel, Looker, DBT, Azure Blob Storage, Azure Data Lake Storage Gen2, Python, Java, SQL Server, Informatica, IICS, Airflow, Terraform, Kubernetes, Docker, Azure Monitor, Oozie, Git, Azure DevOps.

Role: Azure Data Engineer | May 2020 – July 2023
Client: Cisco, Plano, TX
Responsibilities:

- Defined data SLAs and implemented automated data quality checks for Snowflake pipelines, reducing reporting discrepancies and improving stakeholder trust in regulatory and financial analytics.
- Optimized Snowflake warehouse sizing and query execution patterns, reducing compute costs while maintaining performance for high-concurrency analytical workloads.
- Partnered with product, analytics, and compliance teams to translate business and regulatory requirements into Snowflake data models and pipelines, enabling timely delivery of high-impact financial and regulatory reporting.
- Created and managed Snowflake tables tailored to specific data storage and processing needs, using Azure Databricks and Star schema for advanced data modeling and transformation tasks.
- Worked extensively with Snow SQL and Snowpipe, converting Talend Joblets to support Snowflake functionality, resulting in a 20% increase in data processing efficiency. Utilized Python, Java, PySpark, and SQL for custom data processing tasks, enhancing flexibility and performance, and implemented API endpoints for seamless integration and interaction.
- Optimized Power BI reports and dashboards for enhanced performance, ensuring fast loading times and a smooth user experience. This involved optimizing DAX queries, reducing data refresh times, and refining data models.
- Tracked and restored historical data for auditing and analysis using Snowflake's time travel features, contributing to a comprehensive data governance strategy. Utilized Looker for generating reports based on Snowflake connections, with data processed and prepared using Python and PySpark within Azure Databricks.
- Implemented Snowflake caching mechanisms to optimize query performance, resulting in a 35% improvement in data processing efficiency. Additionally, performed data migration of multi-state level data from SQL Server to Snowflake using Python and Snow SQL, orchestrated through Azure Data Factory.
- Applied advanced partitioning techniques in Snowflake, achieving a 30% improvement in query performance. Configured and managed multi-cluster warehouses to enhance system responsiveness, orchestrated through Azure Data Factory. Additionally, I leveraged PySpark and Spark SQL for enhanced data processing capabilities.
- Consulted on Snowflake Data Platform Solution Architecture, specializing in design, development, and deployment strategies. Implemented Change Data Capture (CDC) technology in Talend, simplifying data processing and efficient loading of deltas to a Data Warehouse. Utilized Python and Snow SQL for data migration tasks.
- Established Snowflake development standards and best practices for SQL, data modeling, and pipeline design, improving code consistency, maintainability, and onboarding efficiency for engineering teams.
- Owned production support for Snowflake data pipelines, proactively identifying and resolving performance and data issues, and implementing alerting and runbook processes to improve system reliability and uptime.
- Worked with Azure services, validated Looker reports with Azure Synapse Analytics, and optimized ETL workflows using Azure Data Factory, achieving improved data integration processes. Leveraged Kafka for real-time data streaming to enhance data flow efficiency and developed APIs for seamless integration and interaction.
- Integrated Snowflake with Azure Blob Storage and Azure Data Lake Storage Gen2 (ADLS Gen2) to ensure scalable and efficient data storage solutions, incorporating capabilities for managing Slowly Changing Dimensions (SCD) time travel.
- Led data security and compliance initiatives in Snowflake by implementing role-based access, masking policies, and auditing controls to meet regulatory and enterprise security standards.
- Designed and governed schema evolution strategies in Snowflake, ensuring backward compatibility and minimizing downstream impact as source systems evolved across multiple data domains.
- Used Snow pipe for real-time data ingestion into Snowflake, automating data loading processes and reducing manual intervention. Implemented various strategies related to data security, improving system robustness by 20%.
- Implemented end-to-end data lineage and metadata management for Snowflake assets, improving data discoverability and enabling faster impact analysis for schema and pipeline changes.

Environment: Snowflake, Azure Synapse Analytics, Azure Data Factory, Azure Databricks, Kafka, Talend, Snow SQL, Snowpipe, CDC, Snowflake Time Travel, Looker, DBT, Azure Blob Storage, Azure Data Lake Storage Gen2, Python, Java, SQL Server.

Role: Big Data Developer | Mar 2017 – Apr 2020
Client: Comerica Bank, Frisco, TX

Responsibilities:

- Engineered data pipelines utilizing Flume and Sqoop, enabling seamless ingestion of customer behavioral data into HDFS, resulting in a 15% increase in data ingestion efficiency. Utilized SQL for managing data transformations with CTEs, stored procedures, functions, indexes, and views to optimize data processing and querying capabilities
- Implemented Kafka and Spark Streaming for processing streaming data, achieving a 20% reduction in processing time. Leveraged Python and Spark for custom data processing, enhancing flexibility and performance in handling real-time streams. Developed APIs for seamless integration with these pipelines.
- Built and supported high-volume Hadoop pipelines processing large-scale transactional and behavioral datasets, ensuring reliability and performance across batch and streaming workloads.
- Implemented data validation, reconciliation, and error-handling mechanisms within Hadoop and Spark pipelines, improving data accuracy and downstream analytics reliability.
- Optimized Spark and Hive jobs through partitioning, tuning, and resource optimization, significantly reducing execution time and improving cluster efficiency.
- Designed end-to-end streaming architectures using Kafka and Spark Streaming to support near real-time analytics and monitoring use cases.
- Rigorously used Spark-Scala (RDDs, Data Frames, Spark SQL) with Spark-Cassandra-Connector APIs and Azure services like Databricks and Synapse Analytics for tasks such as data migration and business report generation.
- Conducted efficient data extraction from various sources, including MySQL, leveraging Sqoop, resulting in a 25% acceleration of data transfer processes into HDFS. Orchestrated seamless importing and transformation of data through Hive and MapReduce, leading to a 20% enhancement in data quality.
- Orchestrated the implementation of CI/CD pipelines in the Hadoop environment, resulting in a substantial 30% reduction in deployment time. Utilized Jenkins, Git, Ansible, and Azure services for optimized infrastructure cost reduction.
- Developed and optimized data processing workflows to transform and analyze large datasets. Batch processing with tools like Apache Spark or Apache Hadoop MapReduce, as well as real-time processing with technologies like Apache Flink.
- Managed schema evolution and backward compatibility for Hive and HDFS datasets, minimizing downstream impact as upstream data sources changed over time.
- Implemented data access controls and security policies across Hadoop and Hive environments, ensuring secure handling of sensitive financial data.
- Designed fault-tolerant data pipelines in Hadoop and Spark, implementing retry logic and checkpointing to ensure reliable processing of large-scale batch and streaming workloads.
- Worked closely with platform teams to optimize Hadoop cluster resource utilization, improving job scheduling efficiency and reducing contention across concurrent workloads.
- Contributed to early cloud adoption initiatives by integrating Hadoop-based pipelines with Azure services, laying the foundation for future cloud migration and modernization efforts.
- Monitored production Hadoop and Spark pipelines, proactively identifying and resolving failures to maintain SLAs for critical data processing workflows.
- Implemented a data pipeline using Kafka, Spark, Hive, and Apache Airflow for data ingestion, transformation, and analysis, demonstrating expertise in cutting-edge big data technologies.
- Organized HDFS data into standardized zone-based layouts, improving data discoverability, reuse, and consistency across analytics and reporting teams.
- Implemented job dependency management and scheduling strategies using Apache Airflow and Oozie to ensure reliable execution of complex, multi-stage data workflows.
- Utilized various big data analytic tools, including Hive and MapReduce, to perform in-depth analysis on Hadoop clusters. This initiative led to a 35% improvement in overall analytics efficiency.

Environment: Spark (Spark Streaming, Spark SQL, Spark-Scala, RDDs, Data Frames), Python, Kafka, Sqoop, Apache Flume, Apache Cassandra Azure Services (Azure Data Factory, Azure Databricks, Azure SQL Database, Azure Blob Storage), APIs, Jenkins, Git, Ansible, Maven, Apache Flink, Apache Pig, XML, Agile Methodology.

Role: ETL Developer | Jan 2015 – Feb 2017
Client: Mercury Insurance, Bera, California
Responsibilities:

- Responsible for converting business requirements into ETL specifications and implementing them with Informatica, SQL, PySpark, Shell Scripting, Bash, DB, MDX, and PowerShell, adhering to organizational standards.
- Created Informatica transformations/mapplets/mappings/tasks/worklets/workflows to load data from source to stage, stage to dimensions, bridges, facts, summary, and snapshot facts.
- Experienced using SSIS to create packages to transfer data between Oracle DB, PostgreSQL, MS Access, Flat files, JSON, Excel files, and SQL SERVER 2008 R2. I am also proficient in Python, SSRS, SSAS, SSIS, IICS, and Informatica for comprehensive data integration and reporting solutions.
- Responsible for code migration (through Ansible templates for the Informatica and Autosys Objects) from one phase to another until the final deployment. Managed SSRS, SSAS, SSIS, and IICS for efficient data handling and reporting tasks.
- Ensured compliance with healthcare standards such as HIPAA (Health Insurance Portability and Accountability Act) and HL7 (Health Level Seven International), maintaining data privacy and security in healthcare projects.
- Collaborated with MOCS application services and operations team, leading ETL development efforts, including analysis, design, and integration solutions to meet diverse data and reporting needs of internal and external stakeholders.
- Designed and implemented database structures (DDL), optimized data flow through efficient data manipulation (DML), and ensured data security and access control (DCL) in collaboration with stakeholders.
- Leveraged Matillion's rich library of pre-built components and templates with Azure integration to accelerate ETL pipeline development, reducing development time by 30%.
- Utilized Informatica ETL and Apache Airflow to design and implement Slowly Changing Dimension (SCD) mappings for loading data into Azure SQL Data Warehouse, enhancing data warehouse capabilities and automation.
- Ensured the quality and accuracy of data throughout the ETL process. This included implementing data validation checks, error handling mechanisms, and data reconciliation procedures to detect and resolve data quality issues
- Applied expertise in Informatica PowerCenter and Apache Airflow, creating complex Mappings, Mapplets, Sessions, and Workflows, and conducted performance tuning to optimize large-volume projects.
- Engaged actively in Agile Scrum methodology within an onshore-offshore model, engaging in daily standups and sprints, utilizing Visual SourceSafe, and tracking projects in Trello for streamlined project management.
- Demonstrated proficiency in creating, managing, and delivering reports using SQL Server Reporting Services (SSRS), including the development of interactive web-based reports in Power BI.

Environment: Informatica, Unix Shell Scripting, Bash Scripts, DB Scripts, MDX scripts, PowerShell, SSIS, Python, SSRS, SSAS, IICS, HIPAA, HL7, Matillion, Azure SQL Data Warehouse, DDL, DML, DCL, Agile Scrum, Visual SourceSafe, Trello.